

Anarchy in the Database

Survey and Evaluation of Database Management System Extensibility

Abigale Kim

University of Wisconsin–Madison

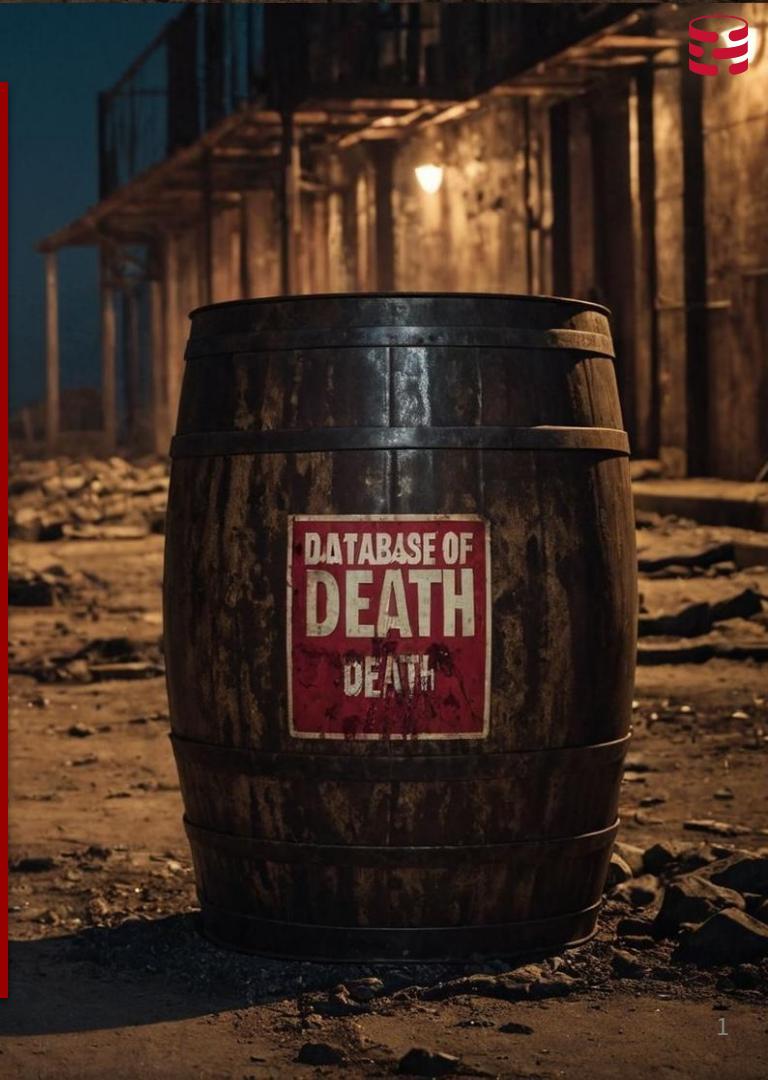
abigale@cs.wisc.edu

Collaborators

Andy Pavlo (pavlo@cs.cmu.edu)

Dave Andersen (dga@cs.cmu.edu)

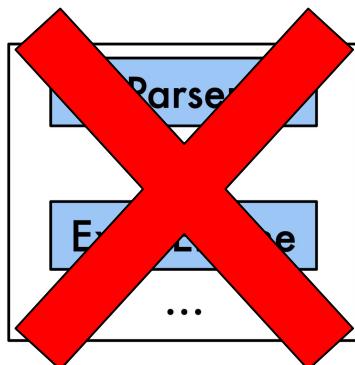
Marco Slot (marco.slot@crunchydata.com)



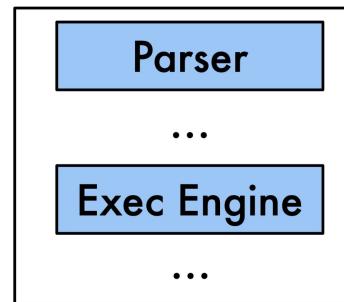


POV: You want to find a DBMS for your use case...

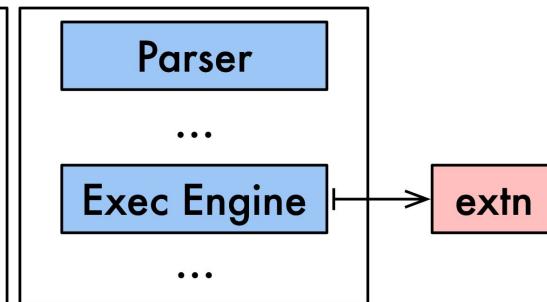
Hmm.. existing DBMSs don't work for our special new use case.



Existing DBMS



Existing DBMS



Existing DBMS
with Extension



What is extensibility?

- Extensibility: the capability of a database system to let custom software extend its capabilities
- Extension: an instance of this software



Extensibility benefits

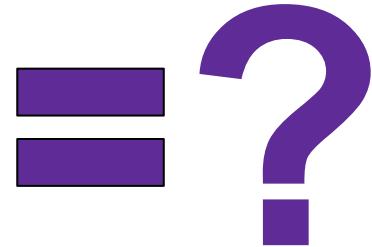
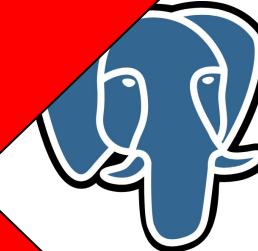
- Support more use cases with less code ✓
- Streamlined developer efforts ✓
- Extensions can be merged into main system ✓
 - autovacuum (merged in 2005)



What happens when we combine extensions?

Distributed PostgreSQL
(as an extension)

auto_explain
(execution logging)



SQL Error [XX000]: ERROR: lookup failed for type 0 with pg auto explain
#7596

Open

StepanYankevych opened this issue 2 weeks ago · 2 comments



Research questions

- How well-designed is current DBMS extensibility?
 - What design decisions caused conflicts like this (and similar)?
 - What can we do to improve the design of DBMS extensibility?
 - Can we design extensibility to not have these conflicts?] future work
-]
- this paper!



Research Layout

Taxonomy and Survey

- Categorize extensibility design and ecosystem (taxonomy)
- Figure out how modern DBMSs support extensibility (survey)

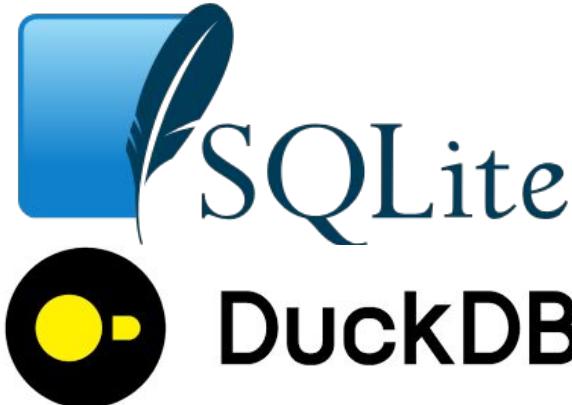
Evaluation

- Collect data on PostgreSQL extensions and analyze the extensibility ecosystem (analysis)



Survey

- Goal: gain larger understanding of the DBMS extensibility design space
- Examined six different DBMSs (PostgreSQL, MySQL, MariaDB, SQLite, DuckDB, Redis OSS)
 - Open source, more comprehensive support for extensibility
- Read extensibility implementation and extensions





Survey Results

	PostgreSQL	MySQL	MariaDB	SQLite	Redis	DuckDB
User-Defined Functions	Yes (408)	Yes (2)	Yes (1)	Yes (79)	Yes (57)	Yes (41)
User-defined Types	Yes (139)	No	Yes (13)	No	No	Yes (4)
Utility Commands	Yes (43)	No	No	No	No	No
Parser Modifications	No				No	Yes (4)
Query Processing	Yes (46)				No	Yes (4)
	Postgres	Too long, can't present: Read the paper!				
Adding Components	Yes				Yes	Yes
Overriding Components	Yes				No	Yes
State Modification	All state				Etxn. + Ephmrl.	All state
Isolation/Security	None				High	Low
Background Workers	Yes				No	No
Memory Allocation	Yes				Yes	Yes
Configuration Options	Yes	Yes	Yes	No	Yes	Yes
Source Code	Yes	Yes	Yes	Yes	No	Yes
Programming Languages	C, C++, Rust	C++	C++	C, Rust	C, Lua	C++
Installation Interface	SQL, configs	SQL	SQL	SQL	SQL, configs	SQL
Build & Test Tooling	Both	Testing	Testing	Both	None	Both
Package Manager	Yes (community)	No	Yes (OS)	Yes (community)	No	Yes



ExtAnalyzer!

Code: <https://github.com/cmu-db/ext-analyzer>

- Tests compatibility between different extensions
 - Compatibility: two extensions work as intended when installed together
- Collects information about extensions
 - Types of extensibility and system components used
 - Source code analysis on bad practices (e.g. duplicate code in extension and DB codebases)

We wrote an analysis framework for all the DBMSs we surveyed, but PostgreSQL has the largest extensibility ecosystem and provided the most fruitful results.



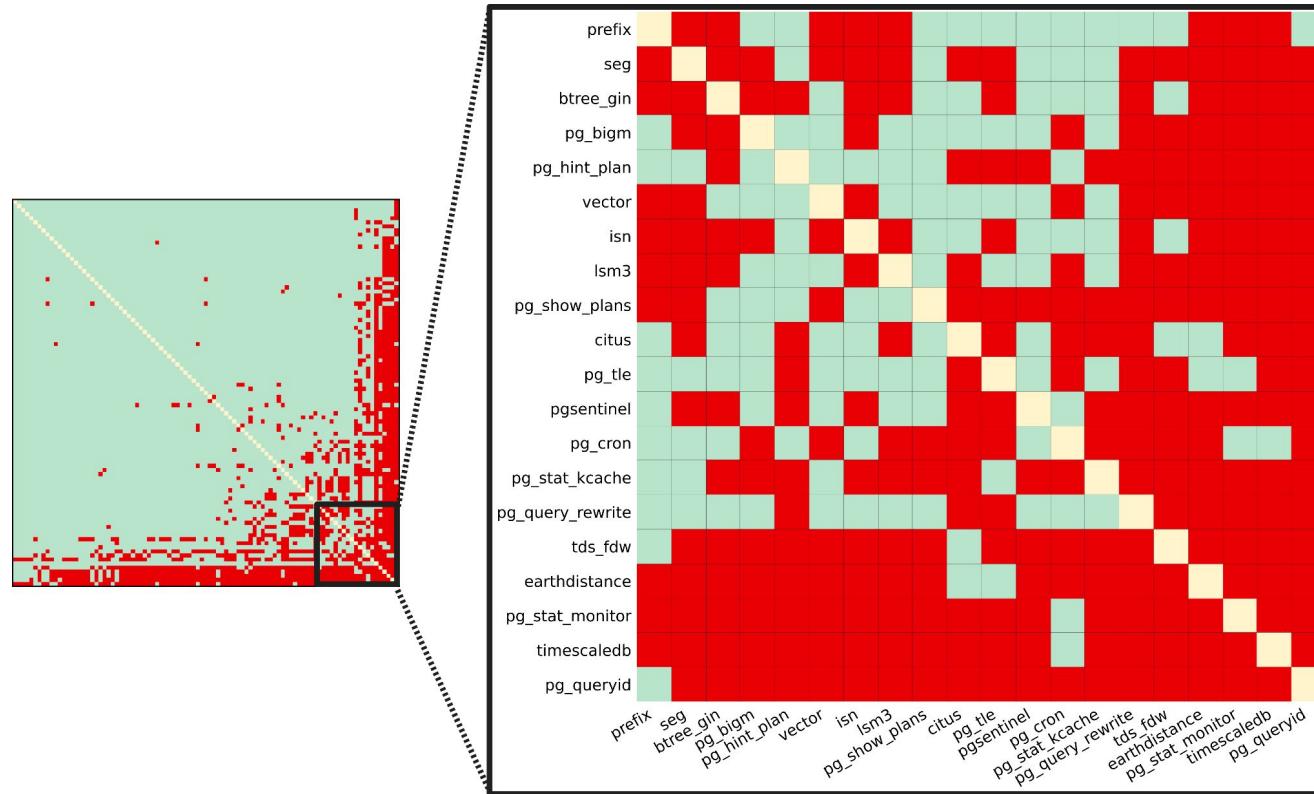
ExtAnalyzer!

- Tests compatibility between different extensions
 - Compatibility: two extensions work as intended when installed together
- Collects information about extensions
 - Types of extensibility and system components used
 - Source code analysis on bad practices (e.g. duplicate code in extension and DB codebases)

We wrote an analysis framework for all the DBMSs we surveyed, but PostgreSQL has the largest extensibility ecosystem and provided the most fruitful results.



Compatibility Failure Matrix





Compatib

PostgreSQL Extens

acl

adminpack

ajbool

argm

pg_hint_plan

https://github.com/osscc-db/pg_hir

The functionality provides a means by allowing users to specify hints via selection of join methods, join order

Information Analysis

Extensibility Types:

System Components:

Duplicate Code Analysis

Total Lines of Code: ?

Lines of Duplicate Code: ?

Percentage of Codebase: ?

Failed Extensions

btree_gin

Terminal output (pg_hint_plan, btree_gin) ▾

Terminal output (btree_gin, pg_hint_plan) ▾

btree_gist

Terminal output (pg_hint_plan, btree_gist) ▲

```
===== installing btree_gist =====
CREATE EXTENSION
===== running regression test queries =====
test init          ... ok    1347 ms
test base_plan     ... ok     15 ms
test pg_hint_plan  ... ok    7437 ms
test ut-init       ... FAILED 831 ms
test ut-A          ... ok    379 ms
test ut-S          ... ok    194 ms
test ut-J          ... ok    133 ms
test ut-L          ... ok     94 ms
test ut-G          ... ok     31 ms
test ut-R          ... ok    444 ms
test ut-fdw        ... ok     37 ms
test ut-W          ... ok    81 ms
test ut-T          ... ok     24 ms
test ut-fini       ... ok     21 ms
test hints_anywhere ... ok     23 ms
test plpgsql       ... ok     80 ms
test oldextversions ... ok    135 ms
```

Terminal output (btree_gist, pg_hint_plan) ▾



Extensibility Distribution





Complex Extensions (>= 3 extensibility types)

- Used K-modes clustering to determine types of complex extensions
- New indexes (67.9%)
 - UDFs, UDTs, index access methods
- Query processing features (21.0%)
 - UDFs, query processing, utility commands
- UDT-optimized query engine (4.9%)
 - UDFs, UDTs, Storage Manager, Query Processing
- New storage manager (4.9%)
 - UDFs, storage manager, utility commands
- Full-featured extensions (1.2%)
 - All extensibility types



Connecting properties to incompatibility

- Ran paired T-tests on extensions and their characteristics
- Higher complexity causes lower incompatibility
 - Using many extensibility types, system components, hooks
 - Duplicate code usage, having a larger codebase
- Larger scope query planner and execution hooks
 - Modifying the planner or executor decreases compatibility



Discussion

- Safety vs. flexibility trade-off
- Extensions naturally compose each other
- Similar extensions exist in different systems
- PostgreSQL API Lessons



Discussion

- Safety vs. flexibility trade-off
- PostgreSQL API Lessons
- Extensions naturally compose each other
- Similar extensions exist in different systems



Safety vs. Flexibility

Low ←← Flexibility →→ High



PostgreSQL



MySQL / MariaDB



DuckDB



SQLite



Redis

Low

↔ Safety ↔

High



PostgreSQL API Lessons

- DBMS should expose finer-grained hooks for common purposes 
- DBMS should have helper APIs to help extensions manage internal state 
- Standardized build and test infrastructure leads to more development 



Takeaways

- Survey findings: PostgreSQL's flexible interface, comprehensive support, and usability results in significantly more prolific ecosystem
- Analysis findings: Extensions are commonly incompatible with each other, caused by higher complexity and query execution modules
- Discussion: safety vs. flexibility, PostgreSQL API gives extensions too much responsibility without helper APIs

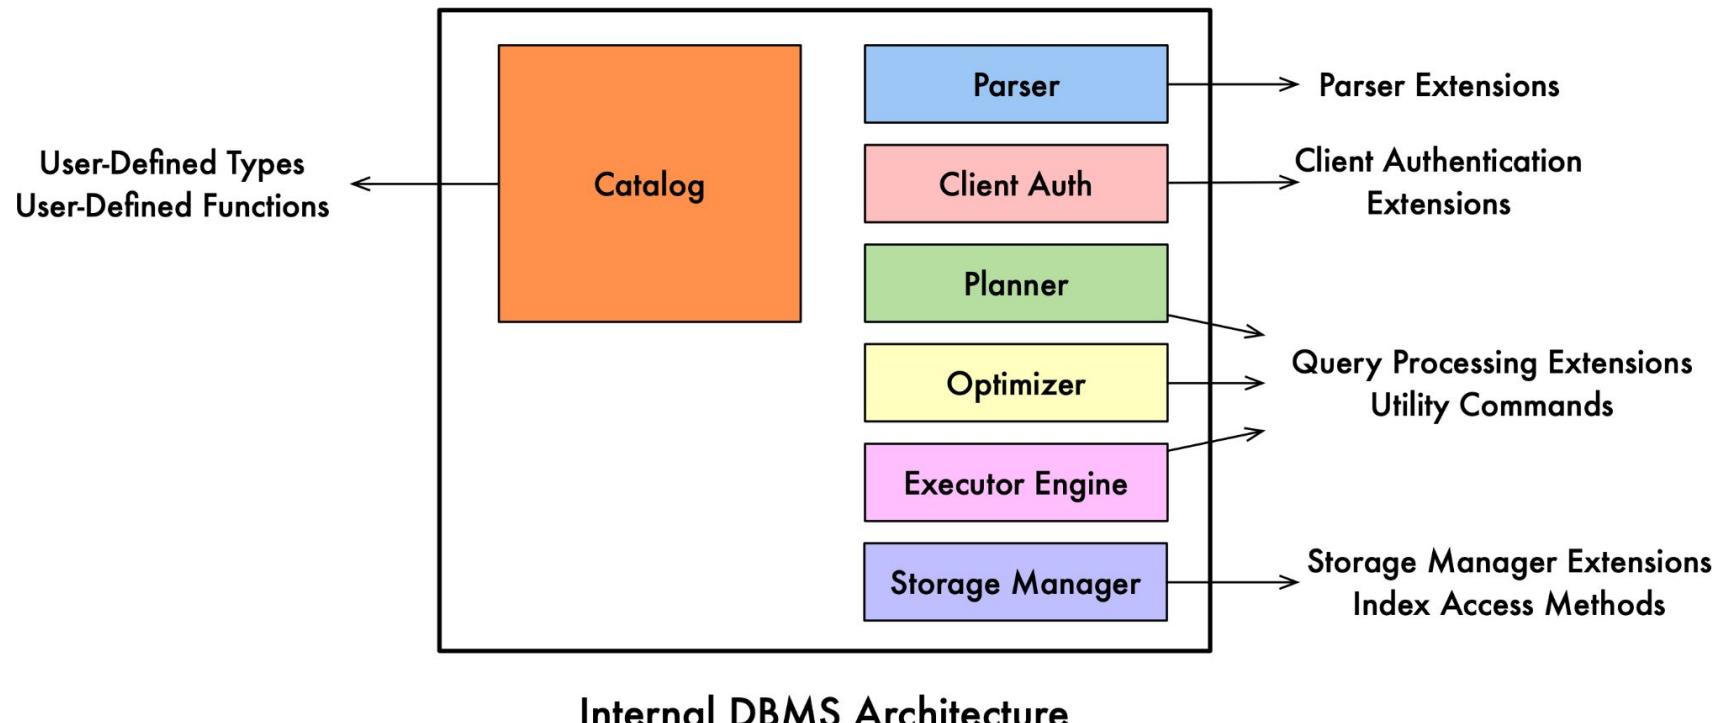
Github Repo: [ExtAnalyzer!](#)

Email: abigale@cs.wisc.edu



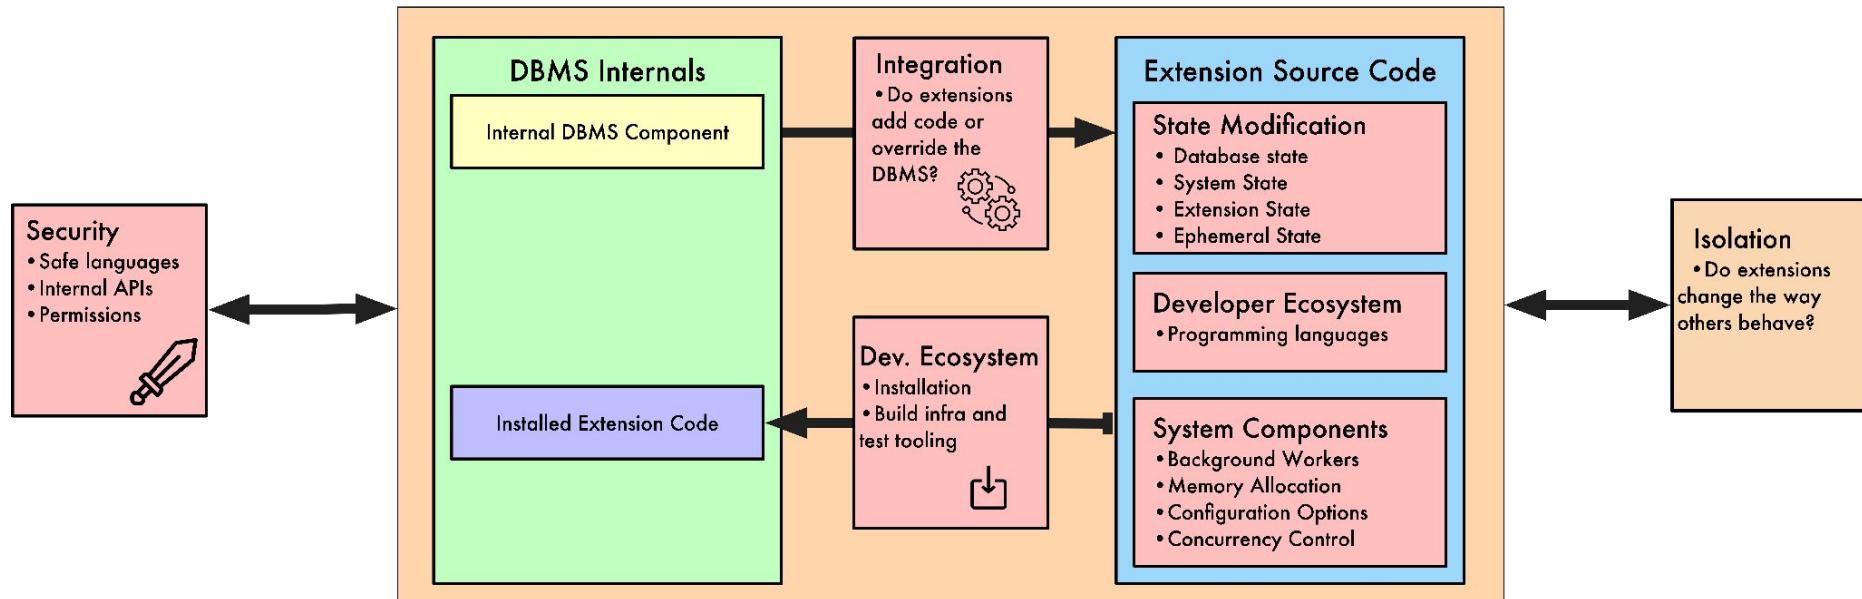


Backup Slide: Types of Extensibility





Backup Slide: Extension Taxonomy Diagram





Backup Slide: Types of Extensibility Survey

	PostgreSQL	MySQL	MariaDB	SQLite	Redis	DuckDB
User-Defined Functions	Yes (408)	Yes (2)	Yes (1)	Yes (79)	Yes (57)	Yes (41)
User-defined Types	Yes (139)	No	Yes (13)	No	No	Yes (4)
Utility Commands	Yes (43)	No	No	No	No	No
Parser Modifications	No	Yes (2)	Yes (1)	No	No	Yes (4)
Query Processing	Yes (46)	Yes (7)	Yes (5)	No	No	Yes (4)
Storage Managers	Yes (44)	Yes (13)	Yes (18)	Yes (43)	No	Yes (9)
Index Access Methods	Yes (67)	No	No	No	No	Yes (3)
Client Authentication	Yes (17)	Yes (3)	Yes (10)	No	No	No
Version Examined	v16	v8	v11	v3	v7	v1
Number of Extensions	441	29	68	98	57	44+



Backup Slide: Survey

	PostgreSQL	MySQL	MariaDB	SQLite	Redis	DuckDB
Adding Components	Yes	Yes	Yes	Yes	Yes	Yes
Overriding Components	Yes	Yes	Yes	Yes	No	Yes
State Modification	All state	All state	All state	All state	Etxn. + Ephmrl.	All state
Isolation/Security	None	Low	Low	Medium	High	Low
Background Workers	Yes	Yes	Yes	No	No	No
Memory Allocation	Yes	Yes	Yes	Yes	Yes	Yes
Configuration Options	Yes	Yes	Yes	No	Yes	Yes
Source Code	Yes	Yes	Yes	Yes	No	Yes
Programming Languages	C, C++, Rust	C++	C++	C, Rust	C, Lua	C++
Installation Interface	SQL, configs	SQL	SQL	SQL	SQL, configs	SQL
Build & Test Tooling	Both	Testing	Testing	Both	None	Both
Package Manager	Yes (community)	No	Yes (OS)	Yes (community)	No	Yes